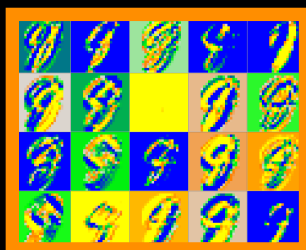




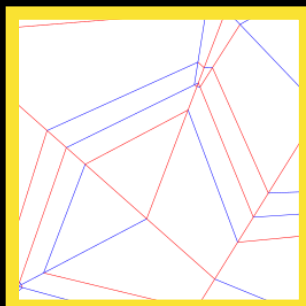
TensorFlow Tutorial, Part 1

Shruti Mishra and Jordan Hoffmann



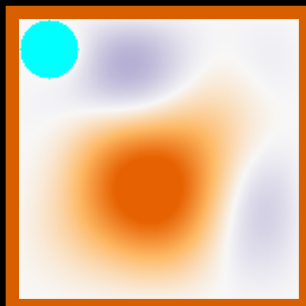
Background

Setup for the aspects we will vary in the coming three demos



Demo 1

Here, we try to predict how many times a sheet has been folded
Classification problem • Geometry • Modify example structure



Demo 2

We solve a PDE and try to predict something about the future
Regression problem • Side Stream • Batch Generator



Demo 3

Given a set of grayscale images, can we recolor them?
Image reconstruction problem • Fully convolutional

IMPORT PACKAGES

- TensorFlow
- Numpy
- Pandas
- Seaborn

DEFINE MODEL

- Different layers
- Sizes, etc

COMPILE MODEL

- Optimizer
- Loss
- Validation

FIT THE MODEL

- How long to train
- When to stop

```
import tensorflow as tf

model = keras.Sequential([
    keras.layers.Flatten(input_shape=(64, 64, 1)),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(10, activation=tf.nn.softmax)
])

model.compile(optimizer=tf.train.AdamOptimizer(),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

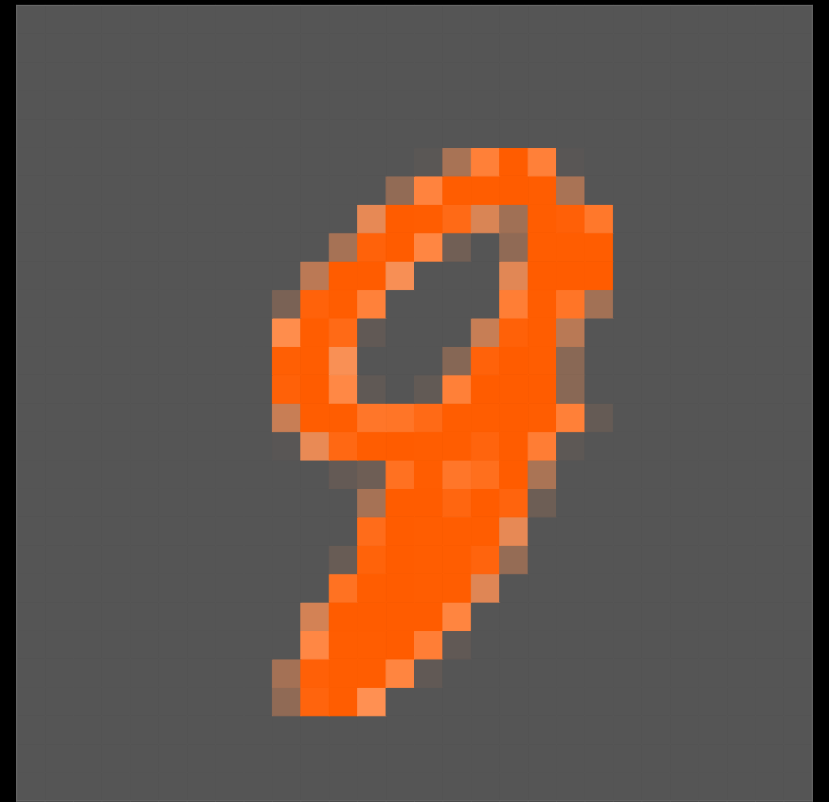
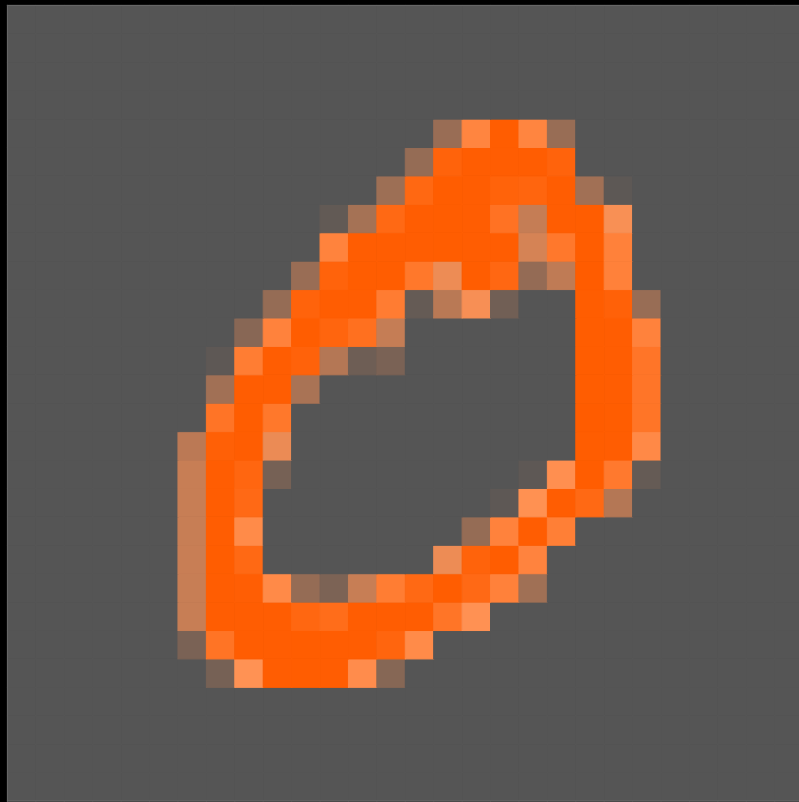
model.fit(x_train, y_train, epochs=15)
```

```
import tensorflow as tf
```

```
model = keras.Sequential([  
    keras.layers.Flatten(input_shape=(64, 64, 1)),  
    keras.layers.Dense(128, activation=tf.nn.relu),  
    keras.layers.Dropout(0.5),  
    keras.layers.Dense(1, activation=tf.nn.relu)  
])
```

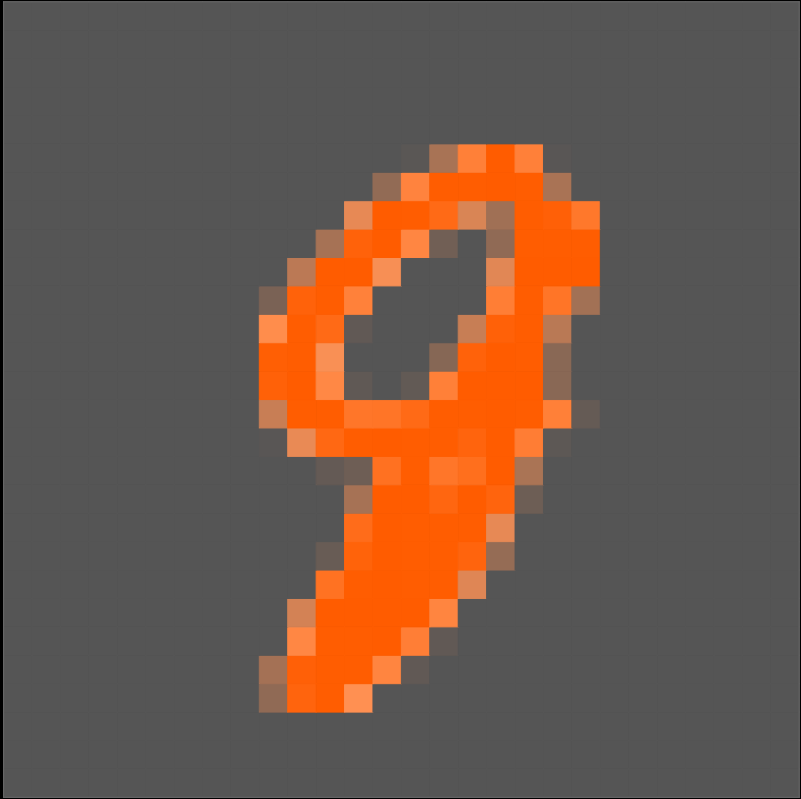
```
model.compile(optimizer='sgd',  
              loss='mean_squared_error')
```

```
model.fit(x_train, y_train, epochs=15, validation_set = (x_test, y_test))
```



Softmax- logistic function

$$\sigma(\vec{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

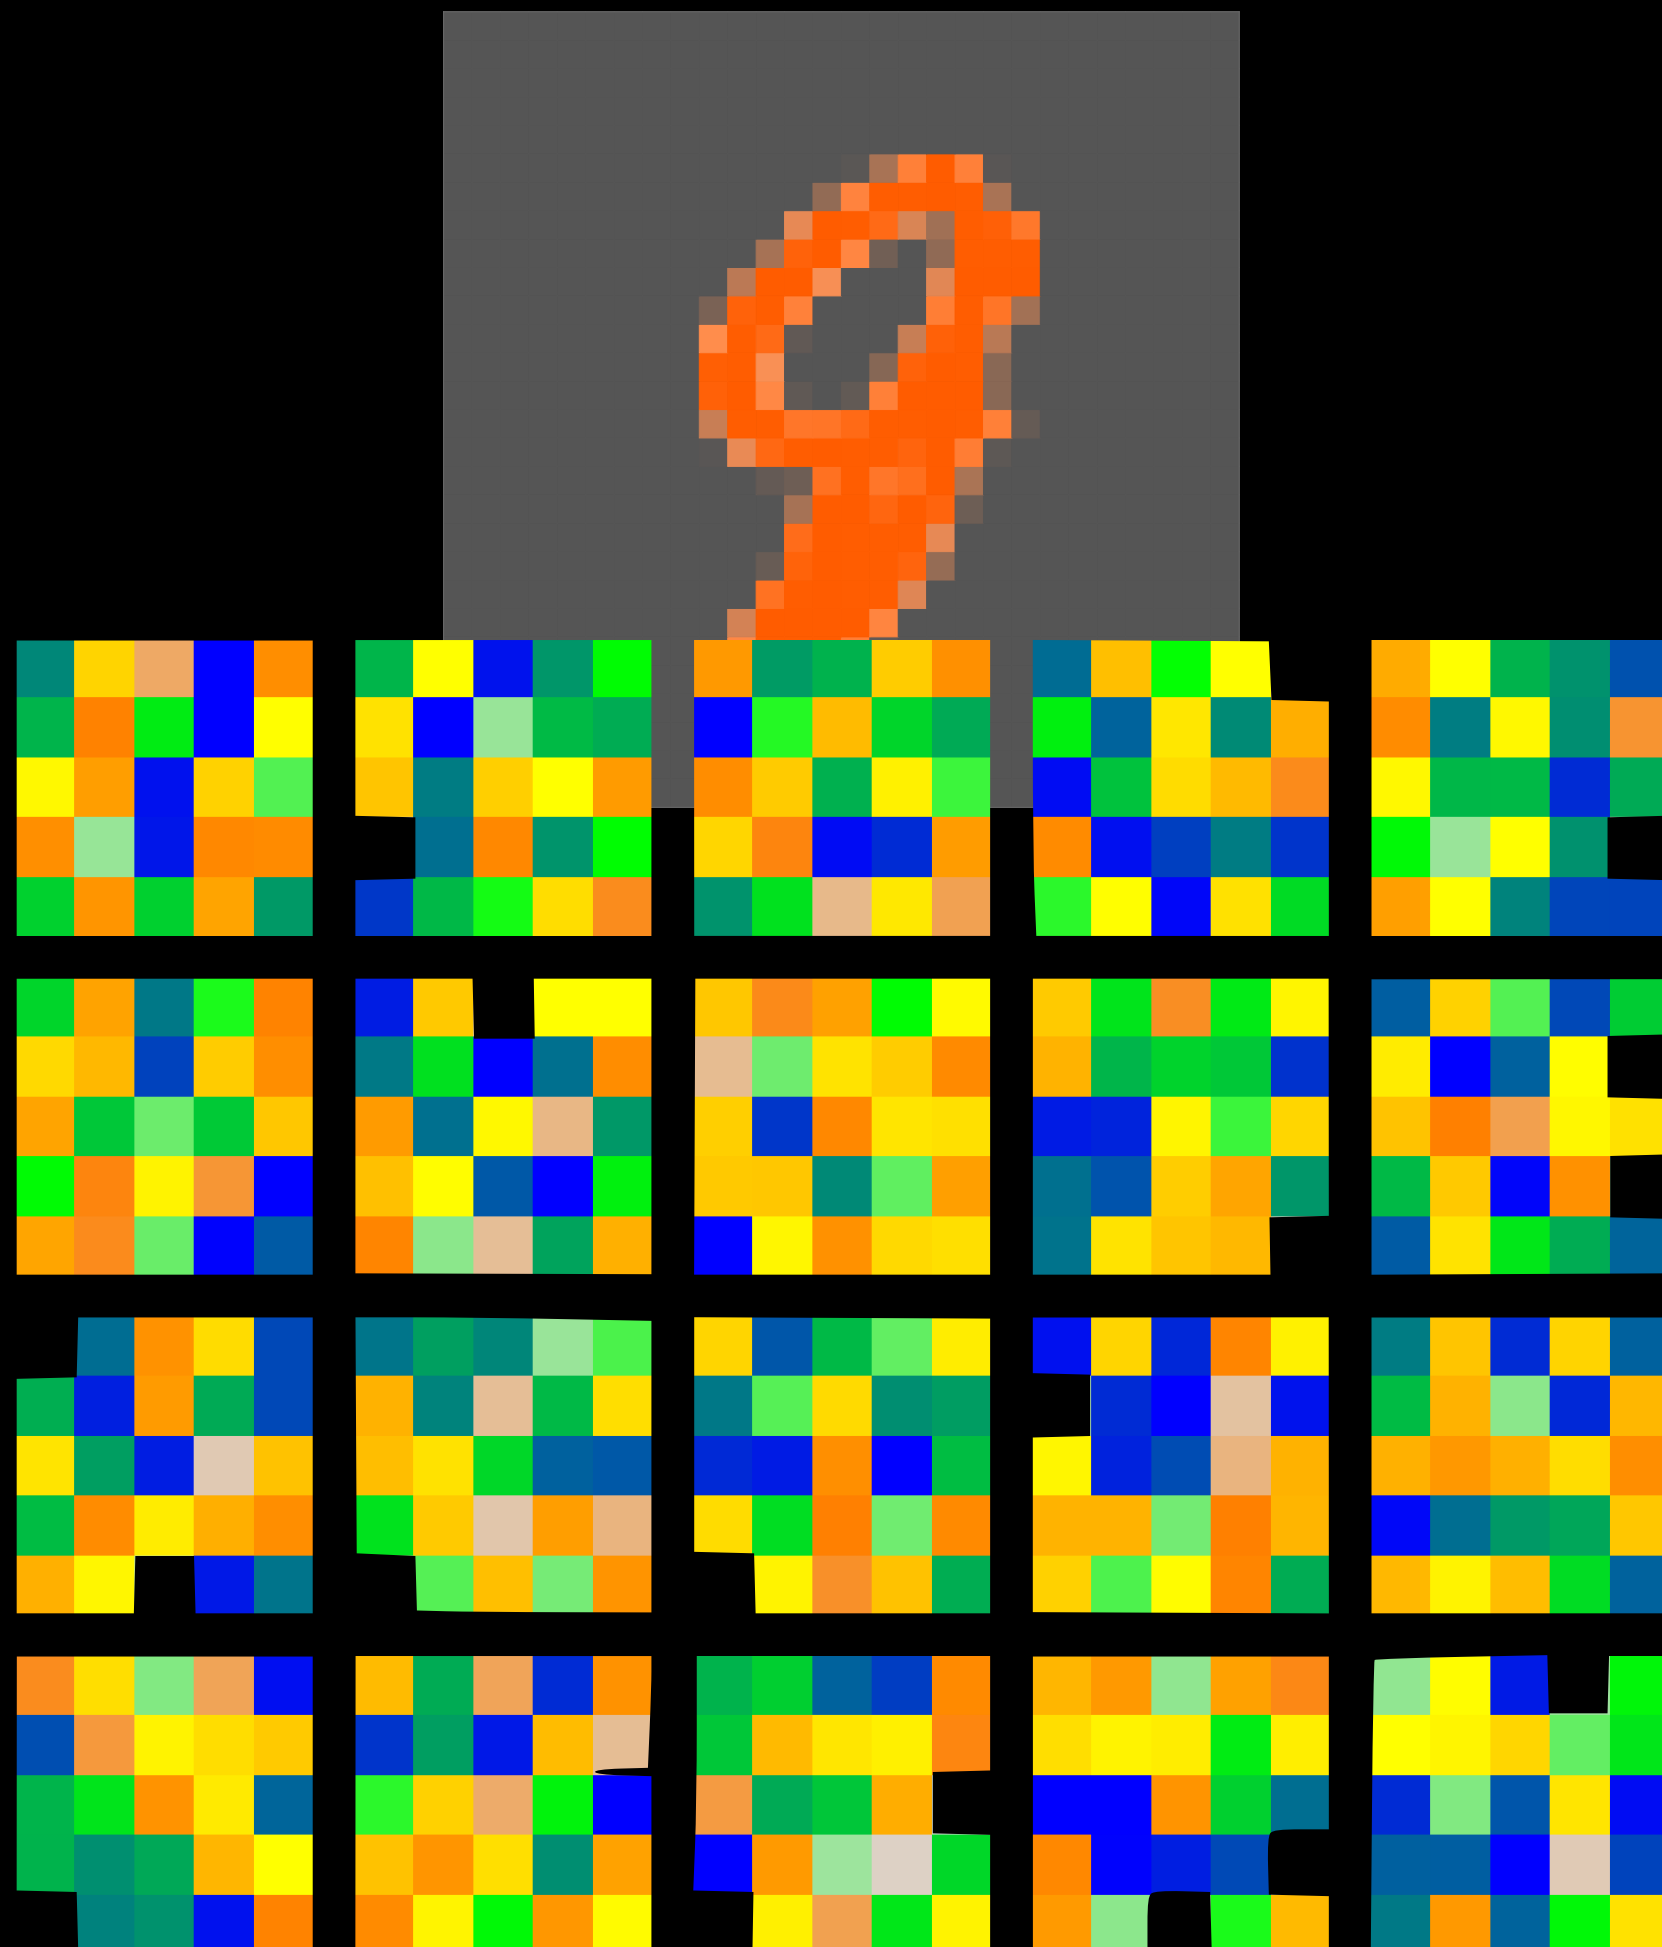




```
import tensorflow as tf

model = keras.Sequential([
    keras.layers.Conv2D(size=(3,3), filters=16, input_shape=(64, 64, 1)),
    keras.layers.MaxPool2D(pool_size=(2,2)),
    keras.layers.Flatten(),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(1, activation=tf.nn.relu)
])

model.compile(optimizer='sgd',
              loss='mean_squared_error')

model.fit(x_train, y_train, epochs=15, validation_set =(x_test,y_test))
```








MaxPool2D

8	9	34	4
3	10	6	5
5	1	8	6
0	4	25	2

10	34
5	25



Mileage distributions of flat folded sheets



1 Fold

2 Folds

3 Folds

4 Folds

5 Folds

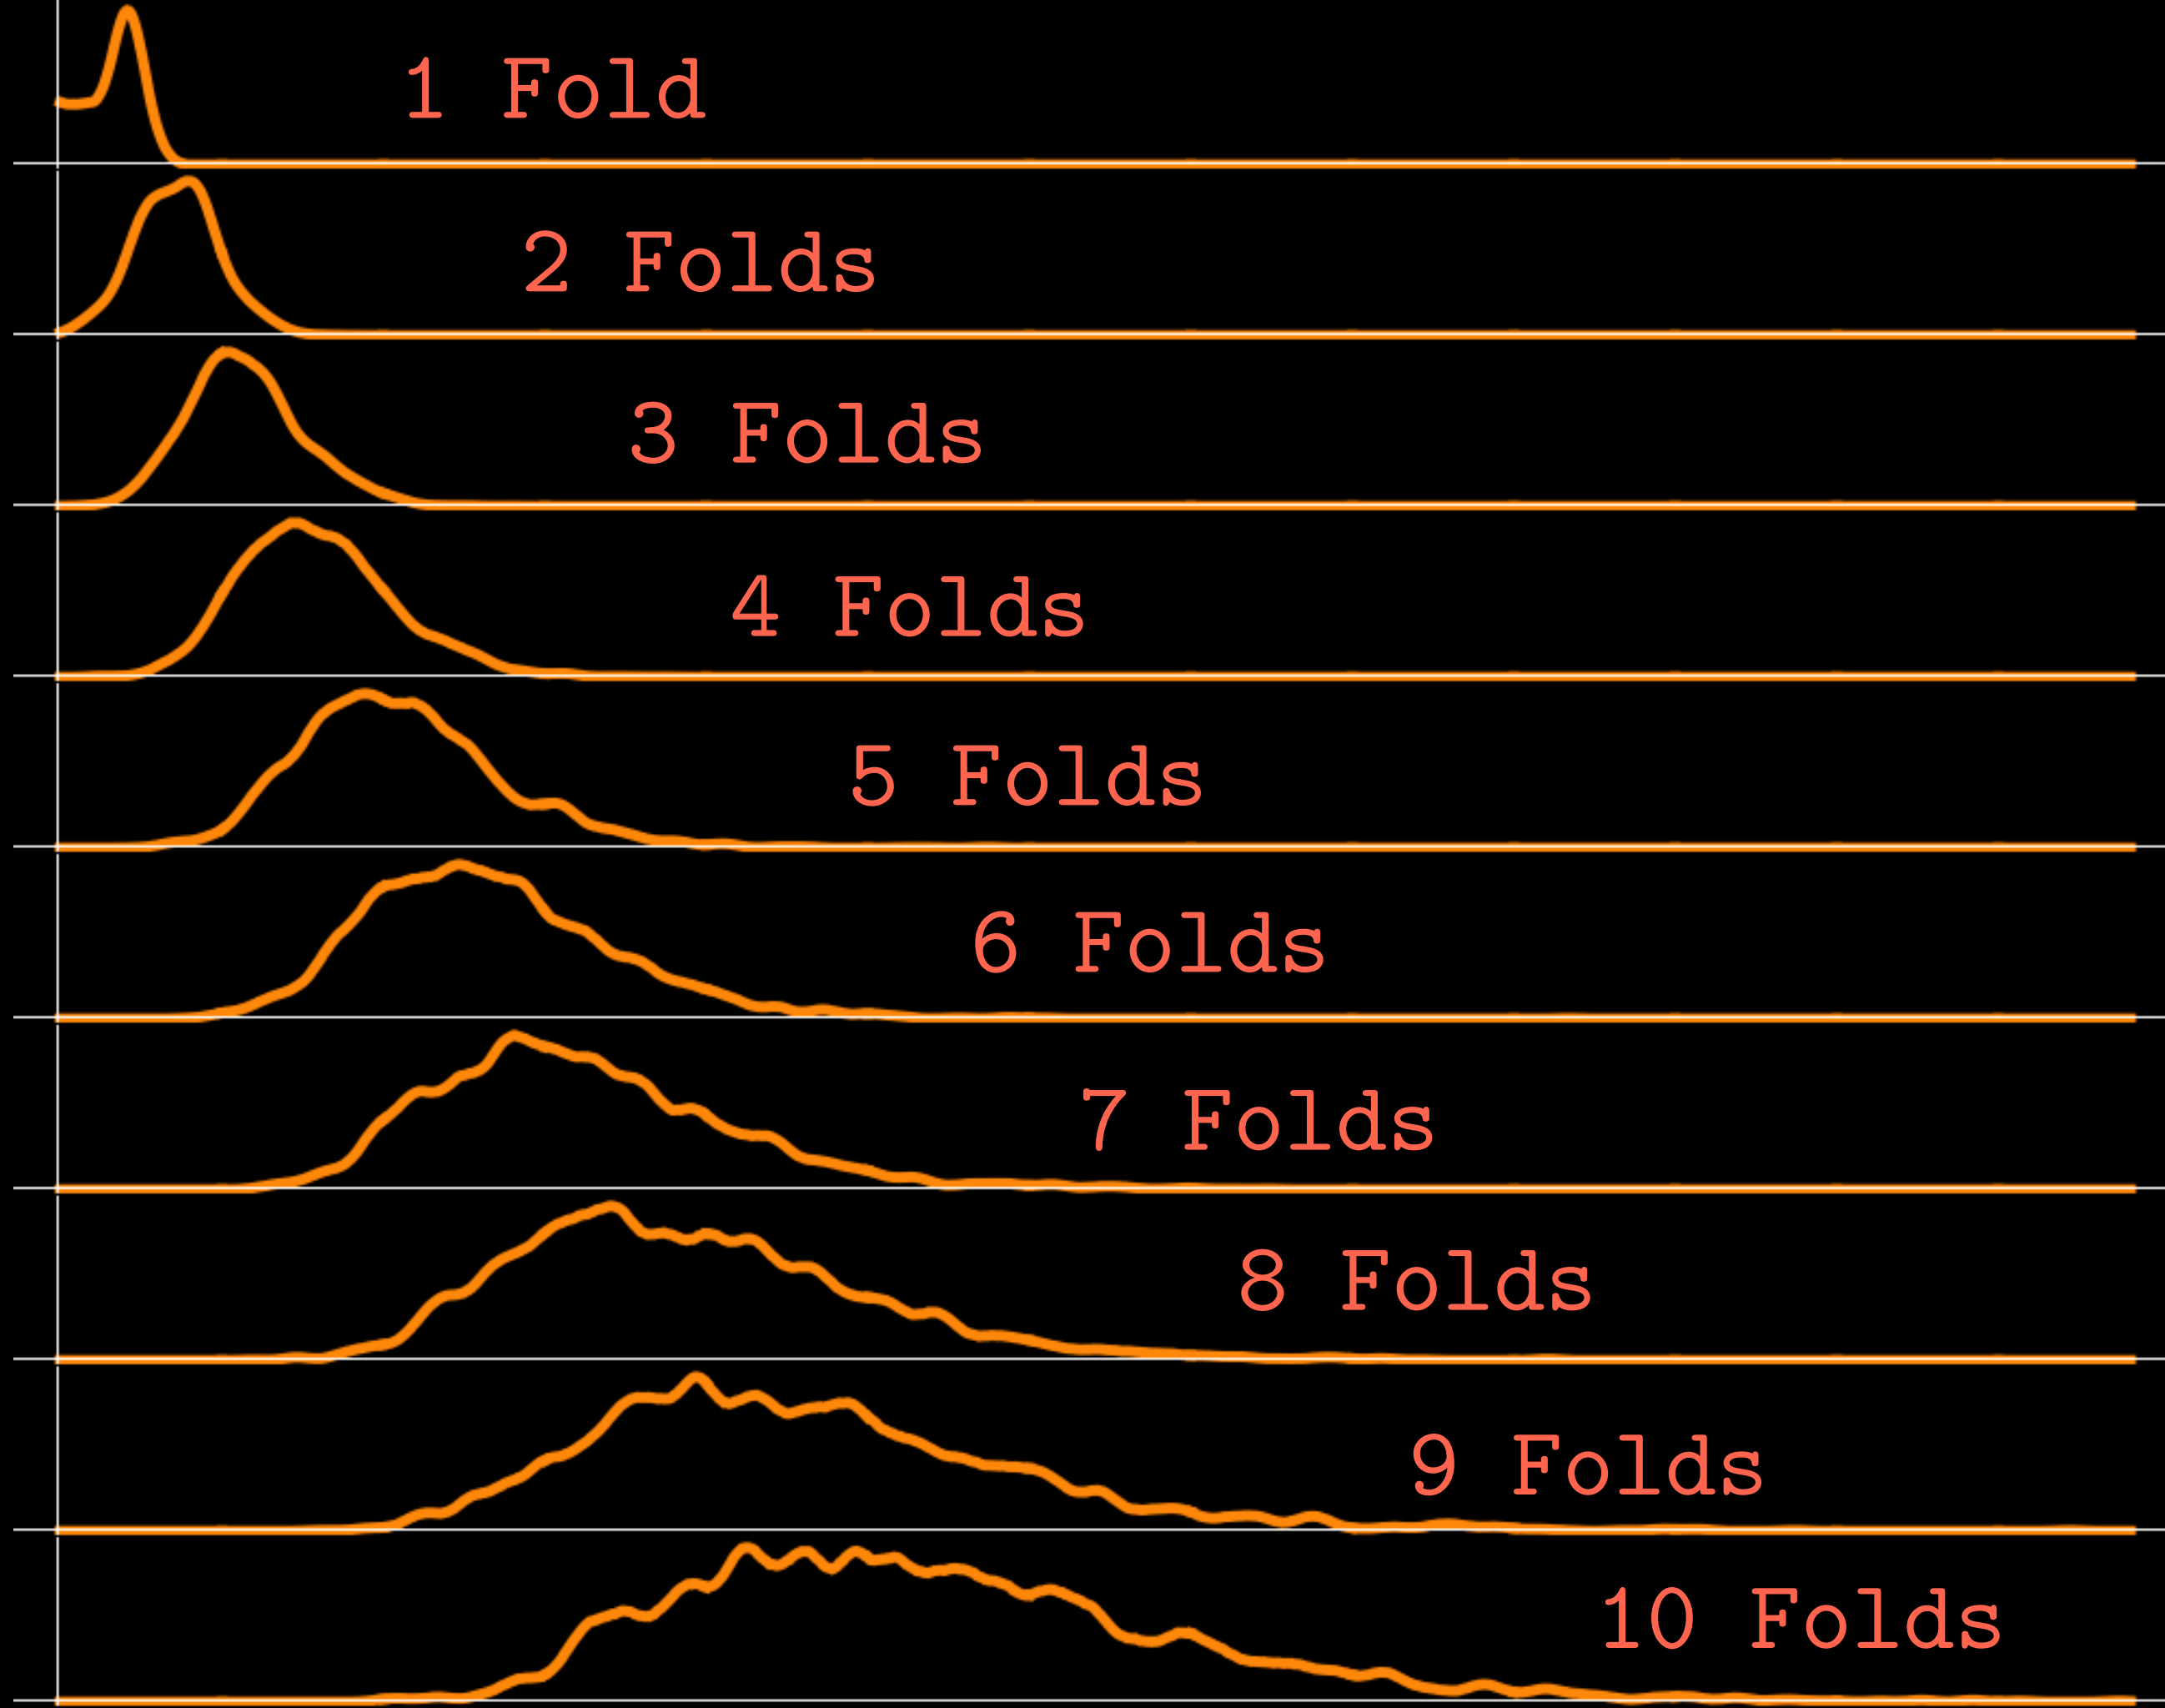
6 Folds

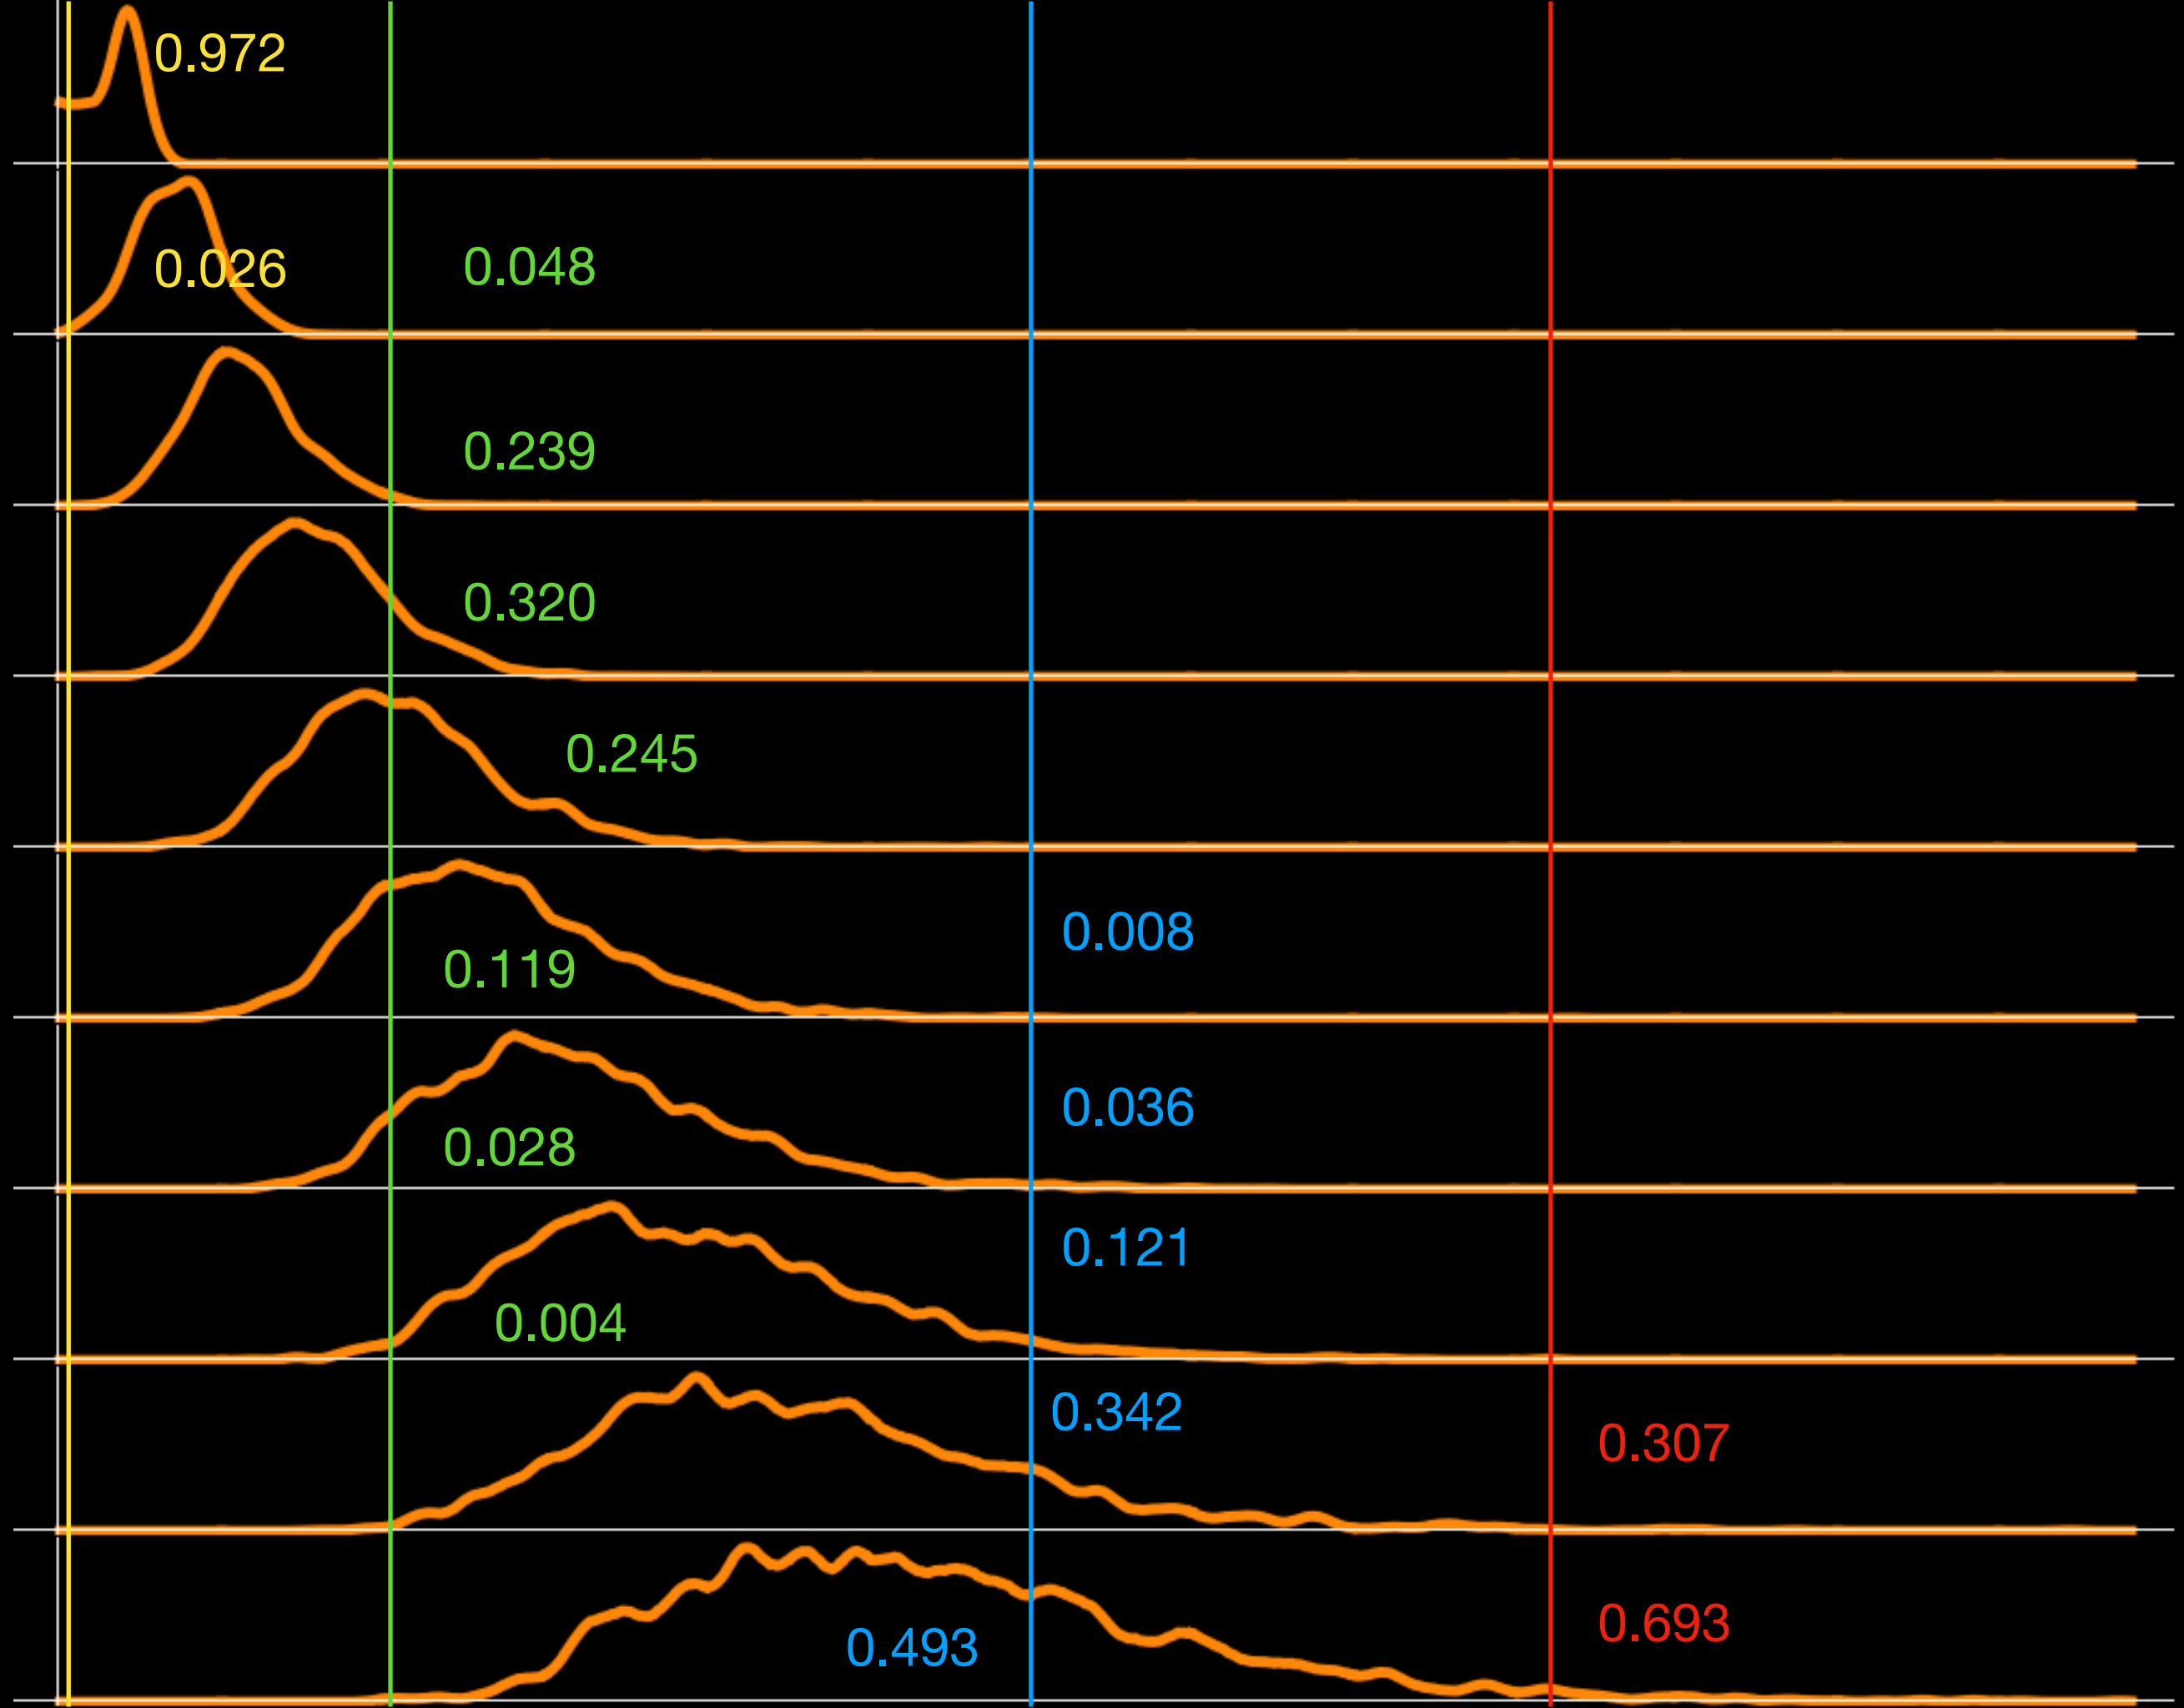
7 Folds

8 Folds

9 Folds

10 Folds





Drawing from these distributions with real creases, we get 32.79% accuracy. Can a computer do better?

